

registry-proxy architecture

Cached, read-only access to agents-repo/registry

Goal

Cloudflare Worker that proxies **only** `agents-repo/registry` files.

Callers (webapp, CLI) avoid hitting GitHub Raw on every request. HTTP is **GET-only** unless an issue explicitly changes method policy. Successful ZIP GETs insert a D1 download event as a side-effect.

Request lifecycle

1. Client hits `https://registry.agents-repo.org` (`/<ref>/<path>` , `/<path>` , `/pkg/...` , `/tags` , or `/stats`)
2. Worker maps to GitHub Raw or Contents API (stats does not proxy)
3. Check `caches.default` by resolved upstream URL
4. On miss, fetch upstream
5. Cache successful 200s and return
6. HTTP 200 versioned ZIPs insert a D1 event in `waitUntil`

Details: `docs/ARCHITECTURE.md` .

`caches.default`

Cloudflare edge cache. Cache-first for repeated GETs.

Catalog files (`index.json` , `tree.json` , `detail.json`) and `/tags` use a **300s** worker TTL. Version snapshot paths stay unbounded at the edge.

Path mapping (no token)

Default upstream:

```
https://raw.githubusercontent.com/agents-repo/registry/<ref>/<path>
```

Example: `/main/packages/index.json` → that file on GitHub Raw.

Omitted ref defaults to `main`.

Path mapping (with token)

When `GITHUB_TOKEN` is present:

GitHub Contents API + `Accept: application/vnd.github.raw`

Better rate-limit behavior. Token comes from **Worker env only**.

GET /tags

Root meta route (not `/main/tags` — that is a file path).

Aggregates GitHub tags API pages into `[{ "name": "v1.2.0" }, ...]`.

Webapp uses this for version-line aliases such as `v2.x`. TTL: 300 seconds.

GET /stats

Root meta route (not `/main/stats` — that is a file path). Prefix like `/pkg/`.

- `/stats` — package totals from D1 `download_events` (`downloads` , `downloads_7d` , `downloads_30d` , `downloads_365d`)
- `?period=all|7d|30d|365d` — ORDER BY only (default `all`); unknown period is 400
- Rolling windows: `datetime('now', '-7 days')` and the same pattern for 30d/365d
- `/stats/packages/<namespace>/<package-id>` — four package totals plus all-time per-artifact rows
- One INSERT per HTTP 200 versioned ZIP GET (cache hit and miss)
- `downloaded_at` is TEXT UTC ISO-8601 from SQLite `datetime('now')`
- Migration `0002` drops undated `download_counts` (intentional reset)
- Missing D1: ZIP still 200; `/stats` is 503

GET-only policy

Supported: GET .

Anything else: 405 Method Not Allowed .

Do not add HTTP write methods without an explicit issue. D1 inserts on ZIP GET 200 are a documented side-effect, not a new method.

Secrets

```
wrangler secret put GITHUB_TOKEN
```

Never commit tokens or put them in `wrangler.jsonc`.

`wrangler secret list` confirms the name without printing the value.

Consumers

- **webapp** production catalog fetch
- **CLI** default org `agents.json` / proxy URL
- Direct curls for debugging (`docs/DEPLOYMENT.md`)

Default catalog URL includes `?ref=v2.x` .

Content-Type

Worker normalizes types from the **path extension** on HTTP 200
(for example `.json` → `application/json`, `.md` → `text/plain`).

Unmapped + `vnd.github.raw` → `application/octet-stream`.

Deploy

Apply D1 migrations first (`wrangler deploy` does not):

```
./scripts/migrate.sh  
./scripts/deploy.sh
```

Then curl `index.json` twice to see cache reuse. See `docs/DEPLOYMENT.md` .

Not this deck

CLI user commands → cli slides.

Package authoring → registry slides.

Links

- `docs/ARCHITECTURE.md`
- `docs/DEPLOYMENT.md`
- `docs/AI_GUIDELINES.md`
- Org ecosystem PDF in `.github` `docs/slides/`

Remember

Read-only HTTP (GET). GitHub Raw or Contents API. Cache hit on repeat GET.
ZIP 200s insert D1 events. Token never in the repo.